
The Human Language
6500
LANGUAGES


The 21st Century
Percentage
Unstructured
Structured
BIG DATA

Text Mining and NLP

Text Mining / Text Analytics is the process of deriving meaningful information from natural language text

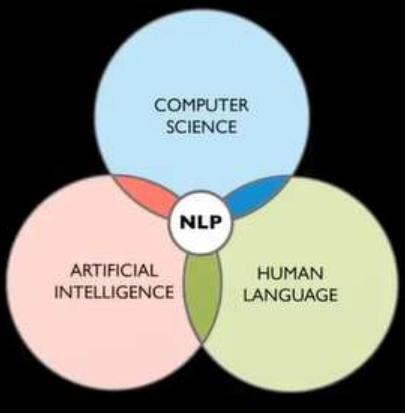

NLP: Natural Language Processing is a part of computer science and artificial intelligence which deals with human languages.

Applications of NLP

Sentimental Analysis

Chatbot

Speech Recognition

Machine Translation

Applications of NLP

Spell Checking

Keyword Searching

Information Extraction

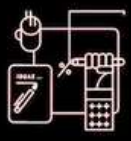

Advertisement Matching

Components of NLP

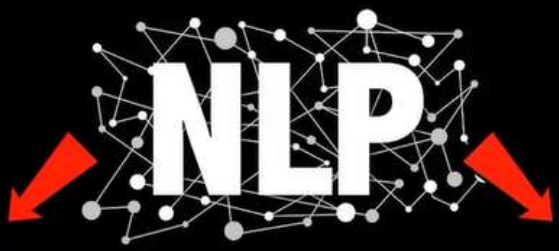

Natural Language Understanding

Natural Language Generation

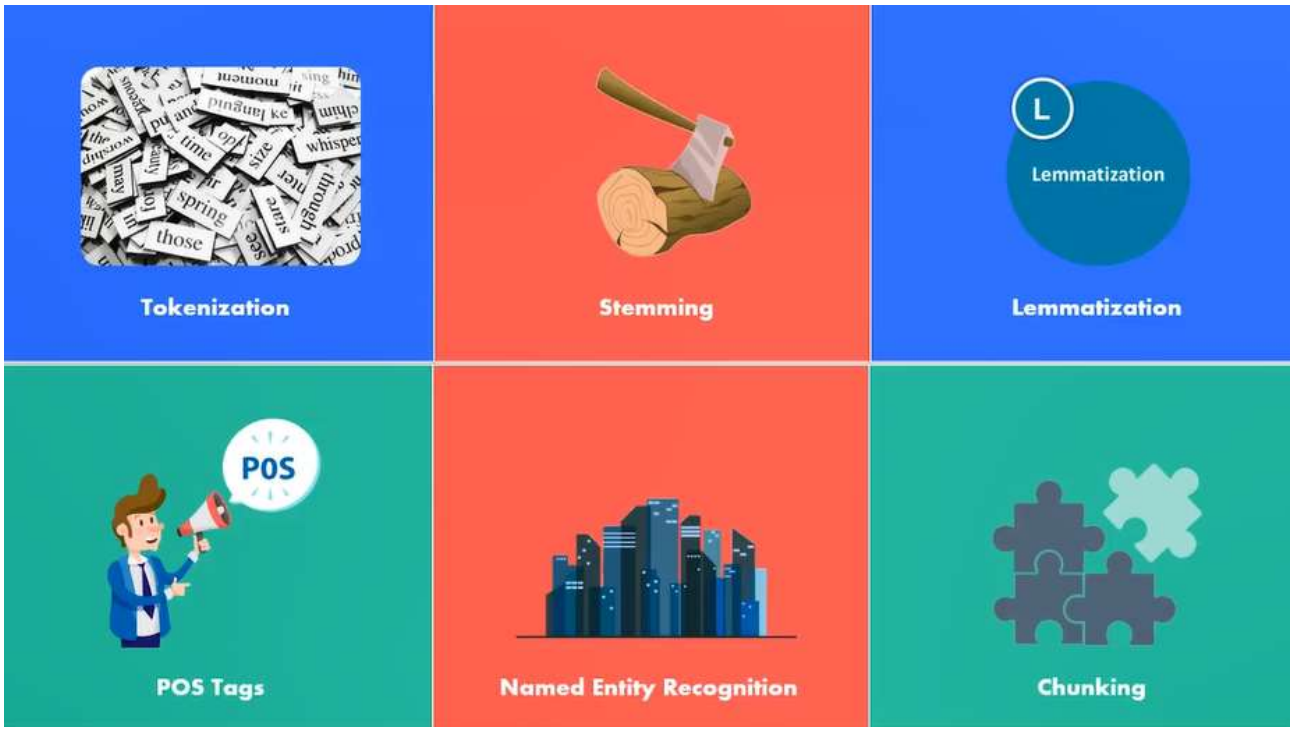

Tokenization

Tokenization is the first step in NLP

Tokenization

Tokenization | is | the | first | step | in | NLP

Stemming

Normalize words into its base form or root form

Affectation | Affects | Affections | Affected | Affection | Affecting

Stemming

Normalize words into its base form or root form

Affect

Lemmatization

Lemmatization

- Groups together different inflected forms of a word, called Lemma
- Somehow similar to Stemming, as it maps several words into one common root
- Output of Lemmatisation is a proper word
- For example, a Lemmatiser should map ***gone***, ***going*** and ***went*** into ***go***

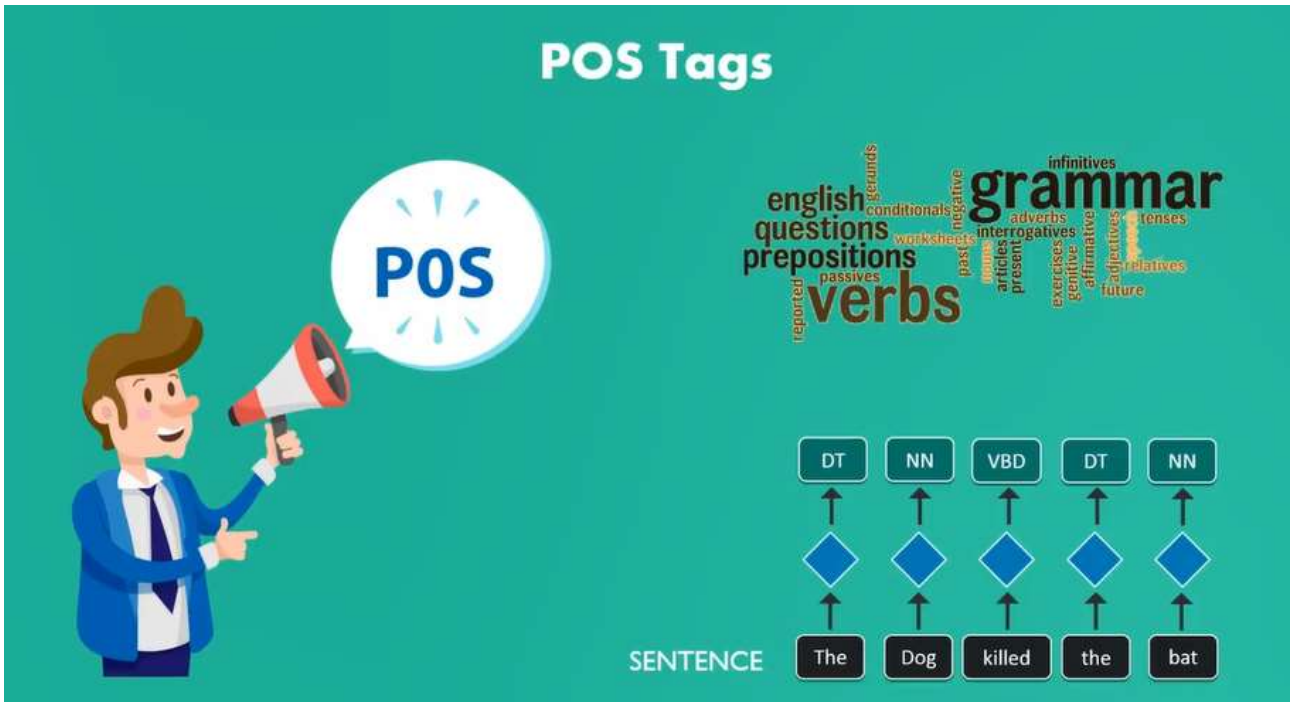
POS Tags
POS
infinitives
grammar
english
gerunds
conditionals
negative
questions
adverbs
tenses
interrogatives
worksheets
prepositions
past
articles
present
exercises
genitive
affirmative
adjectives
relatives
passives
reported
verbs
future
DT
NN
VBD
DT
NN
SENTENCE
The
Dog
killed
the
bat


POS Tags
POS
infinitives
grammar
english
gerunds
conditionals
negative
questions
adverbs
tenses
interrogatives
worksheets
prepositions
past
articles
present
exercises
genitive
affirmative
adjectives
relatives
passives
reported
verbs
future
"Google" something on the internet

Named Entity Recognition

MOVIE

MONETARY VALUE

ORGANIZATION

LOCATION

QUANTITIES

PERSON

Named Entity Recognition

Google's CEO Sundar Pichai introduced the new Pixel3 at New York Central Mall

- Organization
- Person
- Location
- Organization

Chunking

Picking up ***Individual*** pieces of Information and ***Grouping*** them into bigger Pieces

CHUNK

Chunking

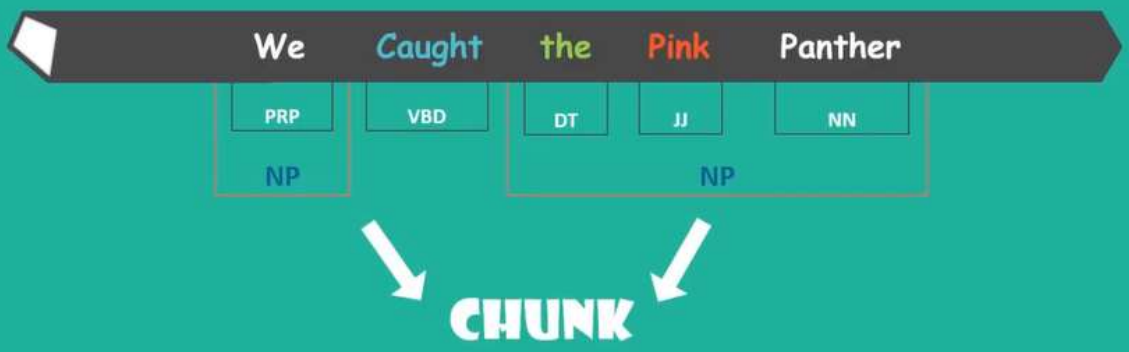

CHUNK

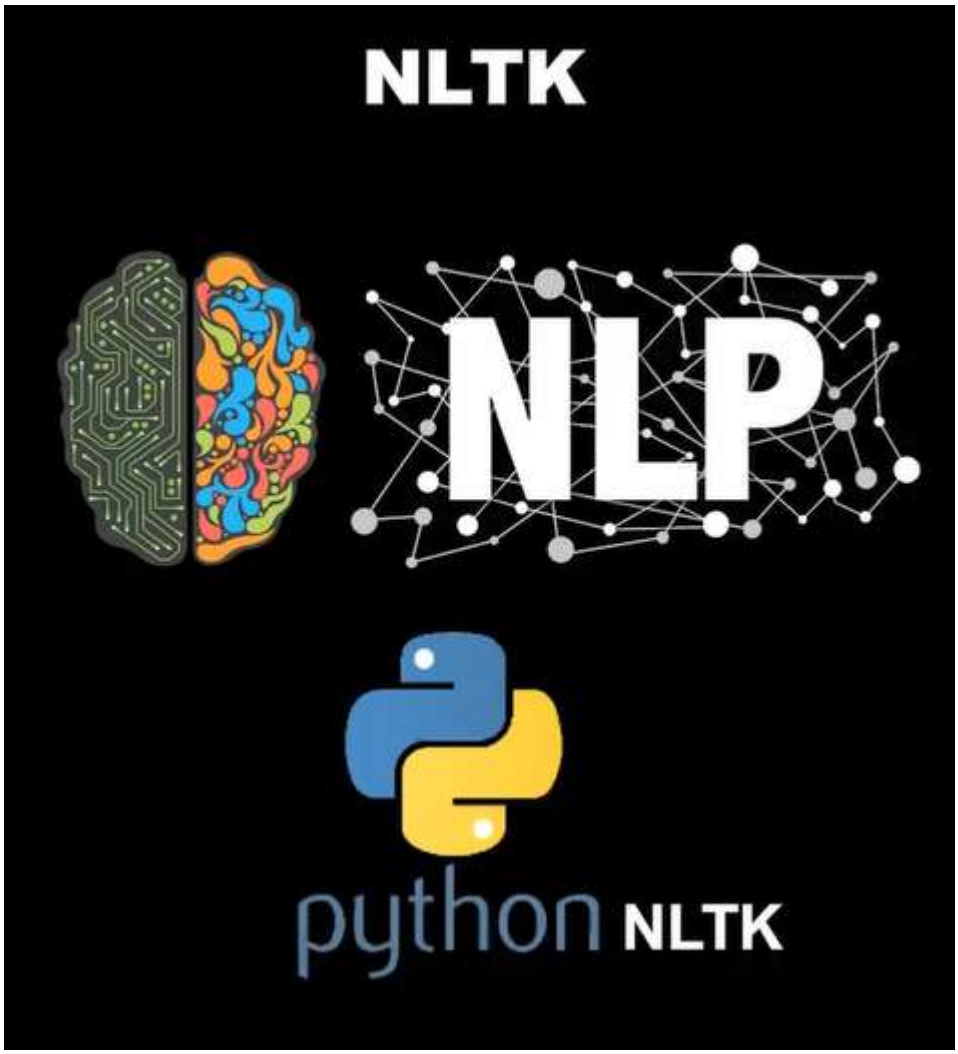

How does spaCy compare to NLTK?

SPACY

- Over 400 times faster
- State-of-the-art accuracy
- Tokenizer maintains alignment
- Powerful, concise API
- Integrated word vectors
- English only (at present)

NLTK

- Slow
- Low accuracy
- Tokens do not align to original string
- Models return lists of strings
- No word vector support
- Multiple languages